

CS221 Section 3: Search

DP, UCS and A*

Motivation

A lot of real-world problems can be modeled as search problems where states have actions that lead to successor states.

- Text segmentation and Vowel insertion
 - Where to put whitespace and vowels, which vowels
- Picking flights to from source city to destination
 - Optimize Time, or Cost, or Layovers
- Structured prediction - e.g. Part-of-Speech, Dependency Parsing
 - Globally Normalized Transition-Based Neural Networks (D. Andor et. al, <https://www.aclweb.org/anthology/P16-1231>)
 - Global Neural CCG Parsing with Optimality Guarantees (K. Lee et. al, <https://arxiv.org/abs/1607.01432>)


Sample Problem

There exists N cities, conveniently labelled from 1 to N .

There are roads connecting some pairs of cities. The road connecting city i and city j takes $c(i,j)$ time to traverse. However, one can only travel from a city with smaller label to a city with larger label (i.e. each road is one-directional).

From city 1 , we want to travel to city N . What is the **shortest time** required to make this trip, given the additional constraint that we should **visit more odd-labeled cities than even labeled cities**?

Example



Best path is [1, 3, 4, 5] with cost 16.

[1, 2, 4, 5] has cost 14 but visits equal number of odd and even cities.

State Representation



Key idea: state

A **state** is a summary of all the past actions sufficient to choose future actions **optimally**.

State Representation

We need to know where we are currently at: **current_city**

We need to know how many odd and even cities we have visited thus far: **#odd, #even**

State Representation: (**current_city, #odd, #even**)

Total number of states: $O(N^3)$

Can We Do Better?

Check if all the information is really required

We store **#odd** and **#even** so that we can check whether **#odd - #even > 0** at **(N, #odd, #even)**

Why not store **#odd - #even** directly instead?

(current_city, #odd - #even) -- $O(N^2)$ states

Precise Formulation of Problem

Let E be the set of one-way roads between cities. $(i, j) \in E$ means there's a road from city i to city j .

State $s = (i, d)$ -- (current_city, #odd - #even)

Start: $(1, 1)$


isGoal(s): $i = N$ and $d > 0$

Actions(s): $\{ \text{move}(j) \text{ for } (i, j) \in E \}$


Cost(s, move(j)): $c(i, j)$

Succ(s, move(j)): $(j, d + 1)$ if j is odd, $(j, d - 1)$ otherwise

Example



State Graph



Solving the Problem

Since we are computing shortest path, which is some form of optimization, we consider DP and UCS.

Recall

- DP can handle negative edges but works only on DAGs
- UCS works on general graphs, but cannot handle negative edges


Since we have a DAG and all edges are positive, both work!

Solving the Problem: Dynamic Programming

$$\text{FutureCost}(s) = \begin{cases} 0 & \text{if IsGoal}(s) \\ \min_{a \in \text{Actions}(s)} [\text{Cost}(s, a) + \text{FutureCost}(\text{Succ}(s, a))] & \text{otherwise} \end{cases}$$


- FutureCost = cost of all future actions
- If state s has no successors, FutureCost(s) is undefined

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	-	-	-
	4	-	-	-	-	-
	5	-	?	-	-	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	-	-	-
	4	-	-	-	-	-
	5	-	?		-	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	-	-	-
	4	-	-	-	-	-
	5	-	?	0	-	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	-	-	-
	4	-	-	-	-	-
	5	-	?	0	0	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	-	-	-
	4	?	-	-	-	-
	5	-	?	0	0	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	-	-	-
	4	?		-	-	-
	5	-	?	0	0	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	-	-	-
	4	?	7	-	-	-
	5	-	?	0	0	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	-	-	-
	4	?	7	7	-	-
	5	-	?	0	0	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	13	-	-
	4	?	7	7	-	-
	5	-	?	0	0	-

Simulation of DP




		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	-	-	-	-
	3	-	-	13	13	-
	4	?	7	7	-	-
	5	-	?	0	0	-

Simulation of DP



		#odd - #even				
		-1	0	1	2	3
city	1	-	-	-	-	-
	2	-	14	-	-	-
	3	-	-	13	13	-
	4	?	7	7	-	-
	5	-	?	0	0	-

Simulation of DP



		#odd - #even				
		-1	0	1	2	3
city	1	-	-	16	-	-
	2	-	14	-	-	-
	3	-	-	13	13	-
	4	?	7	7	-	-
	5	-	?	0	0	-

Solving the Problem: Uniform Cost Search



Algorithm: uniform cost search [Dijkstra, 1956]

Add s_{start} to **frontier** (priority queue)

Repeat until frontier is empty:

 Remove s with smallest priority p from frontier

 If $\text{IsGoal}(s)$: return solution

 Add s to **explored**


 For each action $a \in \text{Actions}(s)$:

 Get successor $s' \leftarrow \text{Succ}(s, a)$

 If s' already in explored: continue

 Update **frontier** with s' and priority $p + \text{Cost}(s, a)$


Simulation of UCS



Explored:
(1, 1) : 0

Frontier:
(3, 2) : 3
(2, 0) : 5

Simulation of UCS



Explored:

(1, 1) : 0


(3, 2) : 3

Frontier:

(2, 0) : 5

(4, 1) : 9

Simulation of UCS



Explored:

(1, 1) : 0

(3, 2) : 3

(2, 0) : 5


Frontier:

(3, 1) : 6

(4, -1) : 7

(4, 1) : 9

Simulation of UCS



Explored:

(1, 1) : 0

(3, 2) : 3

(2, 0) : 5

(3, 1) : 6


Frontier:

(4, -1) : 7

(4, 1) : 9

(4, 0) : 12

Simulation of UCS



Explored:

(1, 1) : 0

(3, 2) : 3

(2, 0) : 5

(3, 1) : 6

(4, -1) : 7


Frontier:

(4, 1) : 9

(4, 0) : 12

(5, 0) : 14

Simulation of UCS




Explored:

(1, 1) : 0
(3, 2) : 3
(2, 0) : 5
(3, 1) : 6
(4, -1) : 7
(4, 1) : 9

Frontier:

(4, 0) : 12
(5, 0) : 14
(5, 2) : 16

Simulation of UCS




Explored:

(1, 1) : 0
(3, 2) : 3
(2, 0) : 5
(3, 1) : 6
(4, -1) : 7
(4, 1) : 9
(4, 0) : 12

Frontier:

(5, 0) : 14
(5, 2) : 16
(5, 1) : 19

Simulation of UCS




Explored:

(1, 1) : 0
(3, 2) : 3
(2, 0) : 5
(3, 1) : 6
(4, -1) : 7
(4, 1) : 9
(4, 0) : 12
(5, 0) : 14

Frontier:

(5, 2) : 16
(5, 1) : 19

Simulation of UCS



Explored:

(1, 1) : 0
(3, 2) : 3
(2, 0) : 5
(3, 1) : 6
(4, -1) : 7
(4, 1) : 9
(4, 0) : 12
(5, 0) : 14
(5, 2) : 16

Frontier:

(5, 1) : 19

STOP!

Comparison between DP and UCS


N total states, n of which are closer than goal state

Runtime of DP is $O(N)$


Runtime of UCS is $O(n \log n)$

DP explores $O(N)$ states.

UCS will explore {1, 2, 5} only.
3 will be in the frontier and 4
will be unexplored.




DP cannot handle cycles



Shortest path is [1, 3, 2, 4, 5] with cost 13.


Hard to define subproblems in undirected graphs

UCS cannot handle negative edge weights



Best path is [1,2,3,4,5] with cost of 8, but UCS will output [1,3,4,5] with cost of 13 because 3 is set as 'explored' before 2.

Improve UCS: A* Search



Recap of A* Search

- Modify the cost of edges and run UCS on the new graph
 - $\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s)$
- $h(s)$ is a heuristic that is our estimate of $\text{FutureCost}(s)$
- If $h(s)$ is consistent then the modified edge weights will return min cost path
 - Consistent: $\text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s) \geq 0$
- One can find a good consistent h by performing relaxation
- If c is min cost on original graph, c' is min cost on modified graph, then $c' = c + h(s_{\text{goal}}) - h(s_{\text{start}})$

Relaxation

A good way to come up with a reasonable heuristic is to solve an easier (less constrained) version of the problem

For example, we can remove the constraint that we visit more odd cities than even cities.

$h(s) = h((i, d)) = \text{length of shortest path from city } i \text{ to city } N$

Note on Relaxation

The main point of relaxation is to attain a problem that **can be solved more efficiently**.

In our case, the modified shortest path problem has $O(N)$ states instead of $O(N^2)$ can thus can be solved more efficiently

Checking consistency

- $\text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s) \geq 0$ (Triangle Inequality)
 - Suppose $s = (i, d)$ and $\text{Succ}(s, a) = (j, d')$
 - Note that $h((i, d)) - h((j, d')) \leq c(i, j) = \text{Cost}(s, a)$
- $h((N, d)) = 0$


How to compute h?

We can reverse the direction of all edges, and then perform UCS starting from city N, and our goal state is city 1.

This takes $O(n \log n)$ time, where n is the number of states whose distance to city N is no farther than the distance of city 1 to city N


city	1	2	3	4	5
h	14	9	13	7	0

Original State Graph



$\text{Cost}(s, a)$


Modified State Graph



$\text{Cost}'(s, a)$

$$= \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s)$$


Simulation of UCS (A*)



Explored:
(1, 1) : 0

Frontier:
(2, 0) : 0
(3, 2) : 2

Simulation of UCS (A*)



Explored:

(1, 1) : 0

(2, 0) : 0


Frontier:

(4, -1) : 0

(3, 2) : 2

(3, 1) : 5

Simulation of UCS (A*)



Explored:

(1, 1) : 0

(2, 0) : 0

(4, -1) : 0


Frontier:

(5, 0) : 0

(3, 2) : 2

(3, 1) : 5

Simulation of UCS (A*)




Explored:

(1, 1) : 0
(2, 0) : 0
(4, -1) : 0
(5, 0) : 0

Frontier:

(3, 2) : 2
(3, 1) : 5

Simulation of UCS (A*)



Explored:

(1, 1) : 0

(2, 0) : 0

(4, -1) : 0

(5, 0) : 0


(3, 2) : 2

Frontier:

(4, 1) : 2

(3, 1) : 5

Simulation of UCS (A*)




Explored:

(1, 1) : 0
(2, 0) : 0
(4, -1) : 0
(5, 0) : 0
(3, 2) : 2
(4, 1) : 2

Frontier:

(5, 2) : 2
(3, 1) : 5

Simulation of UCS (A*)



Explored:

(1, 1) : 0
(2, 0) : 0
(4, -1) : 0
(5, 0) : 0
(3, 2) : 2
(4, 1) : 2
(5, 2) : 2

Frontier:

(3, 1) : 5

STOP!

Actual Cost is $2 + h(1) - h(5) = 2 + 14 - 0 = 16$

Comparison of States visited

UCS

Explored:

(1, 1) : 0

(3, 2) : 3

(2, 0) : 5

(3, 1) : 6

(4, -1) : 7

(4, 1) : 9

(4, 0) : 12

(5, 0) : 14

(5, 2) : 16

Frontier:

(5, 1) : 19

UCS(A*)

Explored:

(1, 1) : 0

(2, 0) : 0

(4, -1) : 0

(5, 0) : 0

(3, 2) : 2

(4, 1) : 2

(5, 2) : 2

Frontier:

(3, 1) : 5

Summary

- States Representation/Modelling
 - make state representation as compact as possible, remove unnecessary information
- DP
 - underlying graph cannot have cycles
 - visit all reachable states, but no log overhead
- UCS
 - actions cannot have negative cost
 - visit only a subset of states, log overhead
- A^*
 - ensure that relaxed problem can be solved more efficiently